

1. (10') Since every conflict-serializable schedule is view serializable, why do we emphasize conflict serializability rather than view serializability?
2. (30') The **lost update** anomaly is said to occur if a transaction T_j reads a data item, then another transaction T_k writes the data item (possibly based on a previous read), after which T_j writes the data item. The update performed by T_k has been lost, since the update done by T_j ignored the value written by T_k .
 - a. Give an example of a schedule showing the lost update anomaly.
 - b. Give an example schedule to show that the lost update anomaly is possible with the read committed isolation level.
 - c. Explain why the lost update anomaly is not possible with the repeatable read isolation level.
3. (20') Consider the following two transactions:

T_1 : read(A);
 read(B);
 if $A = 0$ then $B := B + 1$;
 write(B).

T_2 : read(B);
 read(A);
 if $B = 0$ then $A := A + 1$;
 write(A).

Add lock and unlock instructions to transactions T_1 and T_2 so that they observe the two-phase locking protocol. Can the execution of these transactions result in a deadlock?

4. (40') Consider the following sequence of actions S,
 $S: r1(A), r2(A), r3(B), w1(A), r2(C), r2(B), w2(B), w1(C)$
- and answer the following questions:
- a. Is the schedule S view-serializable? If so, provide a view-equivalent serial schedule.
 - b. Is the schedule S conflict-serializable? If so, describe all the conflict-equivalent serial schedules.
 - c. Is the schedule S a 2PL schedule (with exclusive locks)?
 - d. Is the schedule S a 2PL schedule (with shared and exclusive locks)?